

Zillow Zestimate ADS Nutritional Label

Nicholas Lin

April 2023

1 Zillow ADS Preliminary Analysis

```
[1]: # Importing Dependencies
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import sklearn as sk
```

1.1 Load the Datasets

```
[ ]: # Packages for reading csv file into Colaboratory:
!pip install -U -q PyDrive==1.3.1

from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials

# Authenticate and create the PyDrive client.
# Please follow the steps as instructed when you run the following commands.
auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)
```

```
WARNING: Ignoring invalid distribution -ensorflow
(/usr/local/lib/python3.7/dist-packages)
WARNING: Ignoring invalid distribution -ensorflow
(/usr/local/lib/python3.7/dist-packages)
WARNING: Ignoring invalid distribution -ensorflow
(/usr/local/lib/python3.7/dist-packages)
```

We must load the datasets that are needed to train our machine learning algorithms, handle our data, and to make predictions.

It should be noted that these datasets were the ones provided when entering the competition and accepting the Terms and Conditions.

```
[ ]: # train = pd.read_csv('../input/train_2016_v2.csv' ,  
    ↳ parse_dates=["transactiondate"])  
# properties = pd.read_csv('../input/properties_2016.csv')  
# test = pd.read_csv('../input/sample_submission.csv')  
# test= test.rename(columns={'ParcelId': 'parcelid'})    #To make it easier for  
    ↳ merging datasets on same column_id later
```

```
[ ]: # https://drive.google.com/file/d/1U-1rtYTUk4-_n3GWge_uBjgnsJc80cn4/view?  
    ↳ usp=sharing  
fileid = '1U-1rtYTUk4-_n3GWge_uBjgnsJc80cn4'  
filename = 'train_2016_v2.csv'  
downloaded = drive.CreateFile({'id':fileid})  
downloaded.GetContentFile(filename)  
train = pd.read_csv(filename, parse_dates=["transactiondate"])
```

```
[ ]: train.head(3)
```

```
[ ]:      parcelid  logerror transactiondate  
0    11016594      0.0276      2016-01-01  
1    14366692     -0.1684      2016-01-01  
2    12098116     -0.0040      2016-01-01
```

```
[ ]: train.tail(3)
```

```
[ ]:      parcelid  logerror transactiondate  
90272  12995401     -0.2679      2016-12-30  
90273  11402105      0.0602      2016-12-30  
90274  12566293      0.4207      2016-12-30
```

Since this is for a Colab notebook, we have commented out the above variables in order to implement the csv files from Google Drive directly

```
[ ]: # https://drive.google.com/file/d/1L2EvrHoBREGbOYh6OrScd8PU2Bvm6RVo/view?  
    ↳ usp=sharing  
fileid = '1L2EvrHoBREGbOYh6OrScd8PU2Bvm6RVo'  
filename = 'properties_2016.csv'  
downloaded = drive.CreateFile({'id':fileid})  
downloaded.GetContentFile(filename)  
properties = pd.read_csv(filename)
```

```
/usr/local/lib/python3.7/dist-packages/IPython/core/interactiveshell.py:2882:  
DtypeWarning: Columns (22,32,34,49,55) have mixed types.Specify dtype option on  
import or set low_memory=False.
```

```
    exec(code_obj, self.user_global_ns, self.user_ns)
```

```
[ ]: properties.head(3)
```

```
[ ]:   parcelid  airconditioningtypeid  architecturalstyletypeid  basementsqft  \
0    10754147                    NaN                    NaN          NaN
1    10759547                    NaN                    NaN          NaN
2    10843547                    NaN                    NaN          NaN

   bathroomcnt  bedroomcnt  buildingclasstypeid  buildingqualitytypeid  \
0           0.0          0.0                    NaN                    NaN
1           0.0          0.0                    NaN                    NaN
2           0.0          0.0                    NaN                    NaN

   calculatedbathnbr  decktypeid  ...  numberofstories  fireplaceflag  \
0                NaN          NaN  ...              NaN             NaN
1                NaN          NaN  ...              NaN             NaN
2                NaN          NaN  ...              NaN             NaN

   structuretaxvaluedollarcnt  taxvaluedollarcnt  assessmentyear  \
0                        NaN                9.0          2015.0
1                        NaN             27516.0          2015.0
2             650756.0          1413387.0          2015.0

   landtaxvaluedollarcnt  taxamount  taxdelinquencyflag  taxdelinquencyyear  \
0                9.0          NaN                NaN             NaN
1             27516.0          NaN                NaN             NaN
2             762631.0      20800.37                NaN             NaN

   censustractandblock
0                NaN
1                NaN
2                NaN

[3 rows x 58 columns]
```

```
[ ]: # https://drive.google.com/file/d/15D754PtBPHg7e27bKei6E4d1md2Gwao0/view?
      ↳usp=sharing
fileid = '15D754PtBPHg7e27bKei6E4d1md2Gwao0'
filename = 'sample_submission.csv'
downloaded = drive.CreateFile({'id':fileid})
downloaded.GetContentFile(filename)
test = pd.read_csv(filename)
```

```
[ ]: test.head(3)
```

```
[ ]:   ParcelId  201610  201611  201612  201710  201711  201712
0    10754147        0        0        0        0        0        0
1    10759547        0        0        0        0        0        0
```

```
2 10843547      0      0      0      0      0      0
```

```
[ ]: test = test.rename(columns={'ParcelId': 'parcelid'}) #To make it easier for
      ↪merging datasets on same column_id later
```

##Preliminary research on given data

df_train is a merged dataframe consist of original training data(parcelid and logerror only) and property data(home features)

```
[ ]: df_train_res = train.merge(properties, how='left', on='parcelid')
      df_train_res.head(3)
```

```
[ ]:  parcelid  logerror  transactiondate  airconditioningtypeid  \
0  11016594    0.0276      2016-01-01                1.0
1  14366692   -0.1684      2016-01-01                NaN
2  12098116   -0.0040      2016-01-01                1.0

      architecturalstyletypeid  basementsqft  bathroomcnt  bedroomcnt  \
0                        NaN            NaN            2.0            3.0
1                        NaN            NaN            3.5            4.0
2                        NaN            NaN            3.0            2.0

      buildingclasstypeid  buildingqualitytypeid  ...  numberofstories  \
0                        NaN                4.0  ...                NaN
1                        NaN                NaN  ...                NaN
2                        NaN                4.0  ...                NaN

      fireplaceflag  structuretaxvaluedollarcnt  taxvaluedollarcnt  \
0                NaN                122754.0                360170.0
1                NaN                346458.0                585529.0
2                NaN                61994.0                119906.0

      assessmentyear  landtaxvaluedollarcnt  taxamount  taxdelinquencyflag  \
0                2015.0                237416.0        6735.88                NaN
1                2015.0                239071.0       10153.02                NaN
2                2015.0                57912.0        11484.48                NaN

      taxdelinquencyyear  censustractandblock
0                NaN                6.037107e+13
1                NaN                NaN
2                NaN                6.037464e+13

[3 rows x 60 columns]
```

```
[ ]: df_train_res.shape
```

```
[ ]: (90275, 60)
```

data types

```
[ ]: df_train_res.dtypes.head(34)
```

```
[ ]: parcelid                int64
      logerror               float64
      transactiondate        datetime64[ns]
      airconditioningtypeid  float64
      architecturalstyleid   float64
      basementsqft           float64
      bathroomcnt            float64
      bedroomcnt             float64
      buildingclasstypeid    float64
      buildingqualitytypeid   float64
      calculatedbathnbr       float64
      decktypeid             float64
      finishedfloor1squarefeet float64
      calculatedfinishedsquarefeet float64
      finishedsquarefeet12    float64
      finishedsquarefeet13    float64
      finishedsquarefeet15    float64
      finishedsquarefeet50    float64
      finishedsquarefeet6     float64
      fips                   float64
      fireplacecnt           float64
      fullbathcnt            float64
      garagecarcnt           float64
      garagetotalsqft         float64
      hashottuborspa          object
      heatingorsystemtypeid   float64
      latitude               float64
      longitude              float64
      lotsizesquarefeet       float64
      poolcnt                float64
      poolsizesum            float64
      pooltypeid10            float64
      pooltypeid2             float64
      pooltypeid7             float64
      dtype: object
```

```
[ ]: df_train_res["regionidzip"]
```

```
[ ]: 0      96370.0
      1      96962.0
      2      96293.0
      3      96222.0
      4      96961.0
      ...
```

```

90270    96364.0
90271    96327.0
90272    96478.0
90273    96133.0
90274    96244.0
Name: regionidzip, Length: 90275, dtype: float64

```

```
[ ]: df_train_res.dtypes.tail(30)
```

```

[ ]: poolsize      float64
     pooltypeid10  float64
     pooltypeid2   float64
     pooltypeid7   float64
     propertycountylandusecode  object
     propertylandusetypeid      float64
     propertyzoningdesc         object
     rawcensustractandblock     float64
     regionidcity               float64
     regionidcounty             float64
     regionidneighborhood       float64
     regionidzip                float64
     roomcnt                    float64
     storytypeid               float64
     threequarterbathnbr       float64
     typeconstructiontypeid     float64
     unitcnt                    float64
     yardbuildingsqft17        float64
     yardbuildingsqft26        float64
     yearbuilt                  float64
     numberofstories           float64
     fireplaceflag             object
     structuretaxvaluedollarcnt float64
     taxvaluedollarcnt         float64
     assessmentyear            float64
     landtaxvaluedollarcnt     float64
     taxamount                  float64
     taxdelinquencyflag        object
     taxdelinquencyyear        float64
     censustractandblock       float64
dtype: object

```

null data

```

[ ]: graphData = df_train_res.isna().sum(axis=0).reset_index(name="count")
     graphData = graphData.sort_values(by='count')
     graphData.columns=['features', 'count']
     graphData['ratio']=graphData['count']/df_train_res.shape[0]

```

```
graphData.tail(5)
```

```
[ ]:           features  count    ratio
48   yardbuildingsqft26  90180  0.998948
43           storytypeid  90232  0.999524
5      basementsqft      90232  0.999524
15  finishedsquarefeet13  90242  0.999634
8    buildingclasstypeid  90259  0.999823
```

```
[ ]: graphData.head(5)
```

```
[ ]:           features  count    ratio
0           parcelid      0     0.0
27          longitude      0     0.0
35  propertylandusetypeid      0     0.0
37  rawcensustractandblock      0     0.0
39          regionidcounty      0     0.0
```

```
[ ]: empty_features = graphData[graphData['ratio']>0.95]["features"].tolist()
```

```
[ ]: empty_features
```


```
[ ]: ['finishedsquarefeet15',
      'yardbuildingsqft17',
      'hashottuborspa',
      'taxdelinquencyflag',
      'taxdelinquencyyear',
      'pooltypeid2',
      'pooltypeid10',
      'poolsizesum',
      'decktypeid',
      'finishedsquarefeet6',
      'typeconstructiontypeid',
      'architecturalstyletypeid',
      'fireplaceflag',
      'yardbuildingsqft26',
      'storytypeid',
      'basementsqft',
      'finishedsquarefeet13',
      'buildingclasstypeid']
```

```
[ ]: len(empty_features)
```

```
[ ]: 18
```

```
[ ]: # graphData.plot.barh(y='count',x='features').tick_params(axis='y', width=0.3)
     # # failed graph, too crowded
```

```
[ ]: figs, axes = plt.subplots(figsize=(10,20))
      new_index = np.arange(graphData.shape[0])
      axes.barh(new_index, graphData['count'].values)
      plt.yticks(new_index, graphData["features"].values)
      plt.show()
```

From the index, we extract some interesting features to evaluate:

```
logerror;  
transactiondate  
bathroomcnt;  
bedroomcnt;  
PropertyLandUseTypeID;  
yearbuilt;  
taxvaluedollarcnt;
```

full dictionary of indexes can be found at <https://www.kaggle.com/c/zillow-prize-1/data> under zillow_data_dictionary.xlsx

note:

```
regionidcounty;  
regionidcity;  
regionidcity;  
regionidneighborhood;
```

above features don't have proper dictionaries, we only know numbers but not meanings

ex:

```
[ ]: df_train_res["regionidcounty"].head(2)
```

```
[ ]: 0    3101.0  
     1    1286.0  
     Name: regionidcounty, dtype: float64
```

#####logerror:

```
[ ]: print("mean:\t",df_train_res["logerror"].mean())  
     print("median:\t",df_train_res["logerror"].median())  
     print("std:\t",df_train_res["logerror"].std())
```


```
mean:    0.011457219606756682  
median:  0.006  
std:     0.16107883536718484
```

```
[ ]: df_train_res["logerror"].describe()
```

```
[ ]: count    90275.000000  
     mean      0.011457  
     std       0.161079  
     min      -4.605000  
     25%      -0.025300  
     50%       0.006000
```


```
75%          0.039200
max          4.737000
Name: logerror, dtype: float64
```

```
[ ]: plt.scatter(range(df_train_res.index.size), np.sort(df_train_res.logerror.
↪ values),s=0.4)
plt.xlabel('index')
plt.ylabel('logerror')
plt.show()
```



```
[ ]: plt.hist(df_train_res["logerror"])
```

```
[ ]: (array([3.0000e+00, 2.0000e+00, 4.3000e+01, 1.1500e+02, 7.6615e+04,
          1.3271e+04, 1.5800e+02, 5.1000e+01, 1.3000e+01, 4.0000e+00]),
      array([-4.605 , -3.6708, -2.7366, -1.8024, -0.8682,  0.066 ,  1.0002,
          1.9344,  2.8686,  3.8028,  4.737 ]),
      <a list of 10 Patch objects>)
```




####transactiondate

```
[ ]: type(df_train_res['transactiondate'][0])
```

```
[ ]: pandas._libs.tslibs.timestamps.Timestamp
```

```
[ ]: df_train_res["transactiondate"].groupby(df_train_res["transactiondate"].dt.
      ↪month).count().plot(kind="bar", xlabel="months", ylabel='counts')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f402225b410>
```



group by property land use id

```
[ ]: df_train_res.groupby('propertylandusetypeid', as_index=False).size()
```

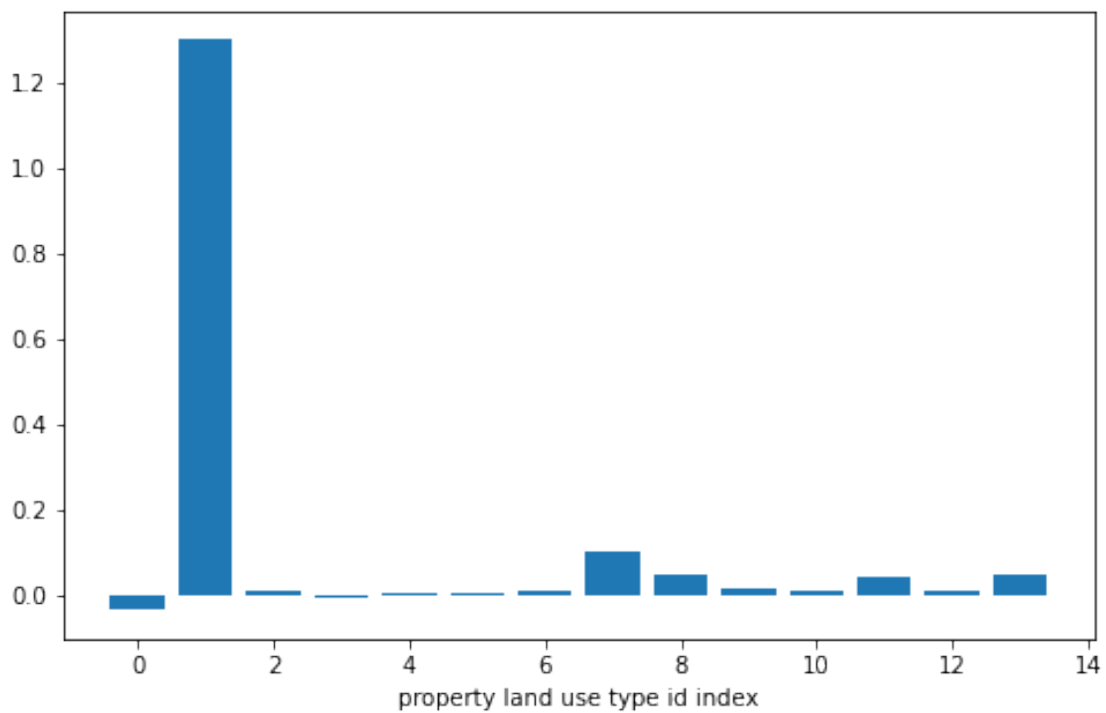
```
[ ]:
  propertylandusetypeid  size
0                31.0     17
1                47.0      1
2               246.0   2376
3               247.0    629
4               248.0    879
5               260.0     62
6               261.0  60637
7               263.0     84
8               264.0     11
9               265.0    356
10              266.0  22815
11              267.0     28
12              269.0   2334
13              275.0     46
```

```
[ ]: df_train_res.groupby('propertylandusetypeid', as_index=False)['logerror'].mean()
```

```
[ ]:
  propertylandusetypeid  logerror
0                31.0 -0.034371
1                47.0  1.301000
```

2	246.0	0.009442
3	247.0	-0.004608
4	248.0	0.005385
5	260.0	0.004606
6	261.0	0.012080
7	263.0	0.104256
8	264.0	0.047591
9	265.0	0.013998
10	266.0	0.010381
11	267.0	0.041118
12	269.0	0.009400
13	275.0	0.050300

```
[ ]: landuse_mean = df_train_res.groupby('propertylandusetypeid',
    ↪as_index=False)['logerror'].mean()
fig, ax = plt.subplots(figsize=(8,5))
plt.bar(range(len(landuse_mean)), landuse_mean["logerror"])
ax.set_xlabel("property land use type id index")
plt.show()
```



####other features

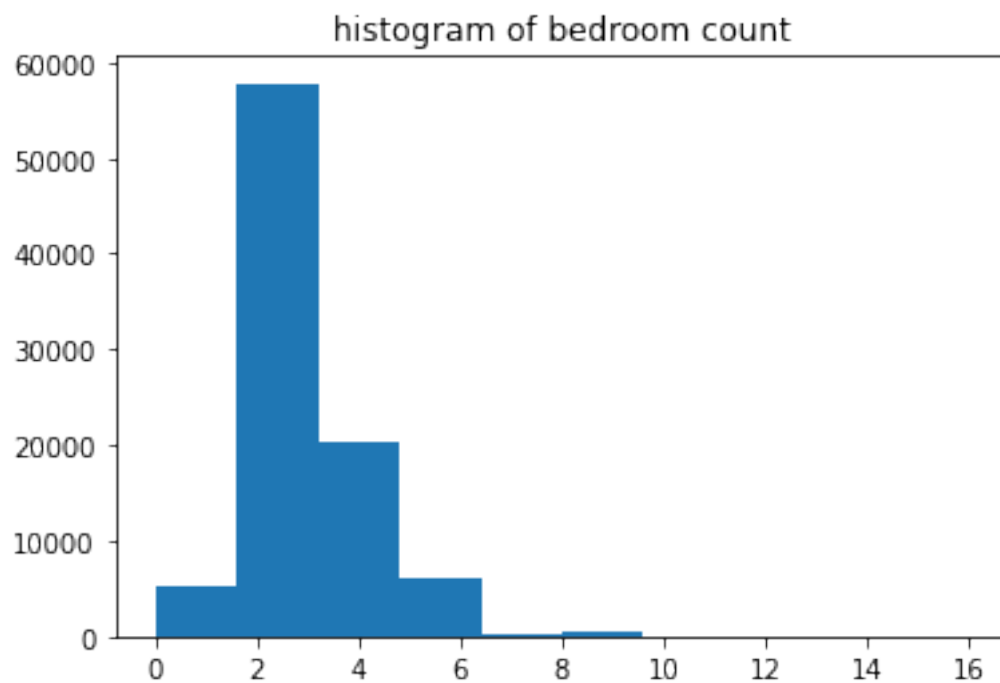
```
[ ]: df_train_res["bedroomcnt"].describe()
```

```
[ ]: count    90275.000000
     mean      3.031869
     std       1.156436
     min       0.000000
     25%       2.000000
     50%       3.000000
     75%       4.000000
     max       16.000000
     Name: bedroomcnt, dtype: float64
```

```
[ ]: df_train_res["bedroomcnt"].median()
```

```
[ ]: 3.0
```

```
[ ]: plt.hist(df_train_res["bedroomcnt"])
     plt.title("histogram of bedroom count")
     plt.show()
```



```
[ ]: df_train_res["bathroomcnt"].describe()
```


```
[ ]: count    90275.000000
     mean      2.279474
     std       1.004271
     min       0.000000
```

```
25%          2.000000
50%          2.000000
75%          3.000000
max          20.000000
Name: bathroomcnt, dtype: float64
```

```
[ ]: df_train_res["bathroomcnt"].median()
```

```
[ ]: 2.0
```

```
[ ]: plt.hist(df_train_res["bathroomcnt"])
plt.title("histogram of bathroom count")
plt.show()
```




```
[ ]: df_train_res["yearbuilt"].describe()
```

```
[ ]: count    89519.000000
mean     1968.532870
std       23.763475
min     1885.000000
25%     1953.000000
50%     1970.000000
75%     1987.000000
max     2015.000000
```


Name: yearbuilt, dtype: float64

```
[ ]: plt.hist(df_train_res["yearbuilt"])  
plt.title("histogram of yearbuilt")  
plt.show()
```



```
[ ]: df_train_res["taxvaluedollarcnt"].describe()
```

```
[ ]: count    9.027400e+04  
mean      4.576726e+05  
std       5.548844e+05  
min       2.200000e+01  
25%      1.990232e+05  
50%      3.428720e+05  
75%      5.405890e+05  
max       2.775000e+07  
Name: taxvaluedollarcnt, dtype: float64
```

```
[ ]: df_train_res["taxvaluedollarcnt"].mean()
```

```
[ ]: 457672.627356714
```

1.1.1 feature correlations

remove columns that have type other than float and int

highly correlated variables are removed from the correlation matrix as they provide repeated and less meaningful information

```
[ ]: train_res_new = df_train_res.drop(empty_features, axis = 1)
```

```
[ ]: train_res_new=train_res_new.select_dtypes('float', 'int')
```

```
[ ]: train_corr=train_res_new.corr().abs()
```

```
[ ]: upper_tri = train_corr.where(np.triu(np.ones(train_corr.shape),k=1).astype(np.  
    ↪bool))
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1:  
DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To  
silence this warning, use `bool` by itself. Doing this will not modify any  
behavior and is safe. If you specifically wanted the numpy scalar type, use  
`np.bool_` here.  
Deprecated in NumPy 1.20; for more details and guidance:  
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations  
    """Entry point for launching an IPython kernel.
```

```
[ ]: to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.  
    ↪95)]
```

```
[ ]: len(to_drop)
```

```
[ ]: 7
```


```
[ ]: train_res_new = train_res_new.drop(to_drop, axis=1)
```

```
[ ]: reduced_corr = train_res_new.corr()
```

```
[ ]: reduced_corr.shape
```

```
[ ]: (31, 31)
```

```
[ ]: fig, ax = plt.subplots(figsize=(10,10))  
    sns.heatmap(reduced_corr)  
    plt.show()
```



2 Original Zillow ADS

Zillow Zestimate ADS: <https://www.kaggle.com/c/zillow-prize-1/overview>

ADS Intended Goal and Specifications: <https://www.kaggle.com/c/zillow-prize-1/data>

Kaggle Original Solution Code: <https://www.kaggle.com/code/zusmani/srpt/script>

```
[ ]: import numpy as np
import pandas as pd
import xgboost as xgb
```

```

import random
import datetime as dt
import gc

import seaborn as sns #python visualization library
color = sns.color_palette()

##matplotlib inline
np.random.seed(1)

```

2.0.1 Analyze the Dimensions of our Datasets

```

[ ]: print("Training Data Size:" + str(train.shape))
      print("Property Data Size:" + str(properties.shape))
      print("Sample Data Size:" + str(test.shape))

```

```

Training Data Size:(90275, 3)
Property Data Size:(2985217, 58)
Sample Data Size:(2985217, 7)

```

2.0.2 Type-Converting the DataSet

Processing some of the algorithms can be made quicker if data representation is made in format **int/float32** instead of the format **int/float64**. Therefore we implement the following lines of code in order to make sure all of our columns are in the float32 type

```

[ ]: for c, dtype in zip(properties.columns, properties.dtypes):
      if dtype == np.float64:
          properties[c] = properties[c].astype(np.float32)
      if dtype == np.int64:
          properties[c] = properties[c].astype(np.int32)

      for column in test.columns:
          if test[column].dtype == int:
              test[column] = test[column].astype(np.int32)
          if test[column].dtype == float:
              test[column] = test[column].astype(np.float32)

```

2.0.3 Let's do some feature engineering

```

[ ]: ### Let's do some feature engineering

      #living area proportions
      properties['living_area_prop'] = properties['calculatedfinishedsquarefeet'] /_
      ↪properties['lotsizesquarefeet']
      #tax value ratio

```

```

properties['value_ratio'] = properties['taxvaluedollarcnt'] /
↳ properties['taxamount']
#tax value proportions
properties['value_prop'] = properties['structuretaxvaluedollarcnt'] /
↳ properties['landtaxvaluedollarcnt']

```

2.0.4 Merging the Datasets

```

[ ]: df_train = train.merge(properties, how='left', on='parcelid')
df_test = test.merge(properties, how='left', on='parcelid')

```

```

[ ]: df_train.head(3)

```

```

[ ]:
  parcelid  logerror  transactiondate  airconditioningtypeid \
0  11016594    0.0276      2016-01-01                1.0
1  14366692   -0.1684      2016-01-01                NaN
2  12098116   -0.0040      2016-01-01                1.0

  architecturalstyletypeid  basementsqft  bathroomcnt  bedroomcnt \
0                      NaN           NaN           2.0           3.0
1                      NaN           NaN           3.5           4.0
2                      NaN           NaN           3.0           2.0

  buildingclasstypeid  buildingqualitytypeid  ...  taxvaluedollarcnt \
0                  NaN                   4.0  ...      360170.0
1                  NaN                   NaN  ...      585529.0
2                  NaN                   4.0  ...      119906.0

  assessmentyear  landtaxvaluedollarcnt    taxamount  taxdelinquencyflag \
0          2015.0          237416.0    6735.879883                NaN
1          2015.0          239071.0   10153.019531                NaN
2          2015.0          57912.0    11484.480469                NaN

  taxdelinquencyyear  censustractandblock  living_area_prop  value_ratio \
0                NaN          6.037107e+13          0.223698    53.470371
1                NaN                NaN          0.621191    57.670429
2                NaN          6.037464e+13          0.194082    10.440699

  value_prop
0    0.517042
1    1.449185
2    1.070486

```

[3 rows x 63 columns]

```

[ ]: df_test.head(3)

```

```
[ ]:  parcelid  201610  201611  201612  201710  201711  201712  \
0  10754147      0      0      0      0      0      0
1  10759547      0      0      0      0      0      0
2  10843547      0      0      0      0      0      0

    airconditioningtypeid  architecturalstyletypeid  basementsqft  ...  \
0                      NaN                      NaN          NaN  ...
1                      NaN                      NaN          NaN  ...
2                      NaN                      NaN          NaN  ...

    taxvaluedollarcnt  assessmentyear  landtaxvaluedollarcnt  taxamount  \
0                9.0          2015.0                9.0          NaN
1            27516.0          2015.0            27516.0          NaN
2        1413387.0          2015.0        762631.0  20800.369141

    taxdelinquencyflag  taxdelinquencyyear  censustractandblock  \
0                NaN                NaN                NaN
1                NaN                NaN                NaN
2                NaN                NaN                NaN

    living_area_prop  value_ratio  value_prop
0                NaN          NaN          NaN
1                NaN          NaN          NaN
2        1.157581    67.950089    0.853304

[3 rows x 67 columns]
```

2.0.5 Remove some unused variables in order to retain some memory

```
[ ]: # del properties, train
# We preserve unused variables for later analysis
gc.collect();

print('Memory usage reduction...')
df_train[['latitude', 'longitude']] /= 1e6
df_test[['latitude', 'longitude']] /= 1e6

df_train['censustractandblock'] /= 1e12
df_test['censustractandblock'] /= 1e12
```

Memory usage reduction...

2.0.6 We will conduct some pre-exploratory analysis to identify the missing values within our datasets.

We incorporated procedures from user Nikunj in order to deal with missing values within our datasets

Reference: <https://www.kaggle.com/nikunjm88/carefully-dealing-with-missing-values>

```
[ ]: # Let's do some engineering with the fireplaceflag variable
print(df_train.fireplaceflag.isnull().sum())
print(df_train.fireplacecnt.isnull().sum())

# By using fireplacecnt variable we can recover some of the fields of
↳fireplaceflag
df_train['fireplaceflag'] = "No"
df_train.loc[df_train['fireplacecnt']>0, 'fireplaceflag'] = "Yes"

# Remaining Missing fireplacecnt will be replaced with 0
index = df_train.fireplacecnt.isnull()
df_train.loc[index, 'fireplacecnt'] = 0

#Tax delinquency flag - assume if it is null then doesn't exist
index = df_train.taxdelinquencyflag.isnull()
df_train.loc[index, 'taxdelinquencyflag'] = "None"
```

90053

80668

```
[ ]: # Similar step performed for Pool/Spa/hot tub
print(df_train.hashottuborspa.value_counts())
print(df_train.pooltypeid10.value_counts())

#lets remove 'pooltypeid10' as has more missing values
print(df_train.hashottuborspa.value_counts())
print(df_train.pooltypeid10.value_counts())

#Assume if the pooltype id is null then pool/hottub doesnt exist
index = df_train.pooltypeid2.isnull()
df_train.loc[index, 'pooltypeid2'] = 0

index = df_train.pooltypeid7.isnull()
df_train.loc[index, 'pooltypeid7'] = 0

index = df_train.poolcnt.isnull()
df_train.loc[index, 'poolcnt'] = 0
```

True 2365

Name: hashottuborspa, dtype: int64

1.0 1161

Name: pooltypeid10, dtype: int64

True 2365

Name: hashottuborspa, dtype: int64

1.0 1161

Name: pooltypeid10, dtype: int64

2.0.7 Label Encoding For Machine Learning & Filling Missing Values

We are now label encoding our datasets. All of the machine learning algorithms employed in scikit learn assume that the data being fed to them is in numerical form. LabelEncoding ensures that all of our categorical variables are in numerical representation. Also note that we are filling the missing values in our dataset with a zero before label encoding them. This is to ensure that label encoder function does not experience any problems while carrying out its operation.

```
[ ]: from sklearn.preprocessing import LabelEncoder

lbl = LabelEncoder()
for c in df_train.columns:
    df_train[c]=df_train[c].fillna(0)
    if df_train[c].dtype == 'object':
        lbl.fit(list(df_train[c].values))
        df_train[c] = lbl.transform(list(df_train[c].values))

for c in df_test.columns:
    df_test[c]=df_test[c].fillna(0)
    if df_test[c].dtype == 'object':
        lbl.fit(list(df_test[c].values))
        df_test[c] = lbl.transform(list(df_test[c].values))
```

2.0.8 Removing the Outliers

```
[ ]: log_errors = df_train['logerror']
df_train = df_train[df_train.logerror < np.percentile(log_errors, 99.5)]
df_train = df_train[df_train.logerror > np.percentile(log_errors, 0.5)]
```

2.0.9 Rearranging the DataSets

We now drop the features which do not serve an useful purpose. We will then split our data and divide it into representation to make it clear which features are to be treated as the determinants in predicting the outcome for our target feature. We make sure to include the same features in the test set as the ones which were included in the training set.

```
[ ]: x_train = df_train.drop(['parcelid', 'logerror', 'transactiondate',
    → 'propertyzoningdesc',
    'propertycountylandusecode' ], axis=1)

x_test = df_test.drop(['parcelid', 'propertyzoningdesc',
    'propertycountylandusecode', '201610', '201611',
    '201612', '201710', '201711', '201712'], axis = 1)

x_train = x_train.values
y_train = df_train['logerror'].values
```


2.0.10 Cross Validation

We divide our datasets into training and validation sets so that we can monitor and test the progress of our machine learning algorithm. This would let us know when our model may be over or under-fitting on a dataset which we have employed.

```
[ ]: from sklearn.model_selection import train_test_split

X = x_train
y = y_train

Xtrain, Xvalid, ytrain, yvalid = train_test_split(X, y, test_size=0.2,
→random_state=42)
```

2.0.11 Implement the Xgboost

We now can select the parameters for Xgboost and monitor the progress of results on our validation set. The explanation of the Xgboost parameters and what they do can be found through the following link: <http://xgboost.readthedocs.io/en/latest/parameter.html>

```
[ ]: dtrain = xgb.DMatrix(Xtrain, label=ytrain)
dvalid = xgb.DMatrix(Xvalid, label=yvalid)
dtest = xgb.DMatrix(x_test.values)

# Try different parameters!
xgb_params = {'min_child_weight': 5, 'eta': 0.035, 'colsample_bytree': 0.5,
→'max_depth': 4,
               'subsample': 0.85, 'lambda': 0.8, 'nthread': -1, 'booster' :
→'gbtree', 'silent': 1, 'gamma' : 0,
               'eval_metric': 'mae', 'objective': 'reg:linear' }

watchlist = [(dtrain, 'train'), (dvalid, 'valid')]

model_xgb = xgb.train(xgb_params, dtrain, 1000, watchlist,
→early_stopping_rounds=100,
                    maximize=False, verbose_eval=10)
```

```
[0]      train-mae:0.474036      valid-mae:0.471898
```

Multiple eval metrics have been passed: 'valid-mae' will be used for early stopping.

Will train until valid-mae hasn't improved in 100 rounds.

```
[10]      train-mae:0.334328      valid-mae:0.332244
[20]      train-mae:0.238427      valid-mae:0.236558
[30]      train-mae:0.173321      valid-mae:0.171623
[40]      train-mae:0.129685      valid-mae:0.128215
[50]      train-mae:0.1011         valid-mae:0.099863
[60]      train-mae:0.082927      valid-mae:0.081904
```

```

[70]    train-mae:0.071828    valid-mae:0.071027
[80]    train-mae:0.065345    valid-mae:0.064763
[90]    train-mae:0.06166     valid-mae:0.061288
[100]   train-mae:0.059599    valid-mae:0.059366
[110]   train-mae:0.058428    valid-mae:0.058296
[120]   train-mae:0.057761    valid-mae:0.057707
[130]   train-mae:0.057369    valid-mae:0.05738
[140]   train-mae:0.057107    valid-mae:0.057194
[150]   train-mae:0.056955    valid-mae:0.057097
[160]   train-mae:0.056837    valid-mae:0.057025
[170]   train-mae:0.056759    valid-mae:0.056985
[180]   train-mae:0.056702    valid-mae:0.056968
[190]   train-mae:0.056649    valid-mae:0.056954
[200]   train-mae:0.056614    valid-mae:0.056949
[210]   train-mae:0.056579    valid-mae:0.056942
[220]   train-mae:0.056542    valid-mae:0.056936
[230]   train-mae:0.056508    valid-mae:0.056931
[240]   train-mae:0.056479    valid-mae:0.056921
[250]   train-mae:0.05645     valid-mae:0.05692
[260]   train-mae:0.056433    valid-mae:0.05693
[270]   train-mae:0.056409    valid-mae:0.056933
[280]   train-mae:0.056384    valid-mae:0.056927
[290]   train-mae:0.05636     valid-mae:0.056927
[300]   train-mae:0.056343    valid-mae:0.056936
[310]   train-mae:0.056315    valid-mae:0.056938
[320]   train-mae:0.056294    valid-mae:0.056938
[330]   train-mae:0.056267    valid-mae:0.056932
[340]   train-mae:0.056246    valid-mae:0.056934
[350]   train-mae:0.056223    valid-mae:0.056931
Stopping. Best iteration:
[251]   train-mae:0.056446    valid-mae:0.056919

```

###Predicting the results###

We can now predict the target variable for our test dataset. All we need to do now is fit the already trained model on the test dataset which we made from merging the sample file with the properties dataset.

```
[ ]: Predicted_test_xgb = model_xgb.predict(dtest)
```

```
[ ]: Predicted_test_xgb
```

```
[ ]: array([-0.05235672, -0.01327515,  0.00347465, ...,  0.10479754,
           0.10479754,  0.10479754], dtype=float32)
```

2.0.12 Submitting the Results

Since we do not need to export the predictions to csv file, we commented out the cells below to save memory

```
[ ]: # # Once again load the file and start submitting the results in each column #
# # sample_file = pd.read_csv('../input/sample_submission.csv')

# # https://drive.google.com/file/d/15D754PtBPHg7e27bKei6E4d1md2Gwao0/view?
→usp=sharing
# fileid = '15D754PtBPHg7e27bKei6E4d1md2Gwao0'
# filename = 'sample_submission.csv'
# downloaded = drive.CreateFile({'id':fileid})
# downloaded.GetContentFile(filename)
# sample_file = pd.read_csv(filename)
# sample_file
```

```
[ ]:
      ParcelId  201610  201611  201612  201710  201711  201712
0      10754147      0      0      0      0      0      0
1      10759547      0      0      0      0      0      0
2      10843547      0      0      0      0      0      0
3      10859147      0      0      0      0      0      0
4      10879947      0      0      0      0      0      0
...          ...    ...    ...    ...    ...    ...
2985212  168176230      0      0      0      0      0      0
2985213   14273630      0      0      0      0      0      0
2985214  168040630      0      0      0      0      0      0
2985215  168040830      0      0      0      0      0      0
2985216  168040430      0      0      0      0      0      0
```

[2985217 rows x 7 columns]

```
[ ]: # for c in sample_file.columns[sample_file.columns != 'ParcelId']:
#     sample_file[c] = Predicted_test_xgb

# print('Preparing the csv file ...')
# sample_file.to_csv('xgb_predicted_results.csv', index=False, float_format='%.
→4f')
# print("Finished writing the file")
# # failed to write to the csv file in drive, but predictions are stored in the
→variable sample_file
```

Preparing the csv file ...
Finished writing the file

```
[ ]: # len(sample_file)
```

```
[ ]: 2985217
```

```
[ ]: # sample_file
```

```
[ ]:      ParcelId    201610    201611    201612    201710    201711    201712
0      10754147 -0.052357 -0.052357 -0.052357 -0.052357 -0.052357 -0.052357
1      10759547 -0.013275 -0.013275 -0.013275 -0.013275 -0.013275 -0.013275
2      10843547  0.003475  0.003475  0.003475  0.003475  0.003475  0.003475
3      10859147  0.116168  0.116168  0.116168  0.116168  0.116168  0.116168
4      10879947  0.061894  0.061894  0.061894  0.061894  0.061894  0.061894
...      ...      ...      ...      ...      ...      ...
2985212 168176230  0.104798  0.104798  0.104798  0.104798  0.104798  0.104798
2985213  14273630  0.104798  0.104798  0.104798  0.104798  0.104798  0.104798
2985214 168040630  0.104798  0.104798  0.104798  0.104798  0.104798  0.104798
2985215 168040830  0.104798  0.104798  0.104798  0.104798  0.104798  0.104798
2985216 168040430  0.104798  0.104798  0.104798  0.104798  0.104798  0.104798
```

[2985217 rows x 7 columns]

3 ADS Evaluation

####Sample submission logerror predictions:


```
[ ]: def describe_arr(array):
      print("mean:",array.mean())
      print("median:",np.median(array))
      print("min:",array.min())
      print("max:",array.max())
```

```
[ ]: describe_arr(Predicted_test_xgb)
```

```
mean: 0.013021984
median: 0.011305124
min: -0.25805414
max: 0.34636483
```

```
[ ]: plt.scatter(range(len(Predicted_test_xgb)), np.sort(Predicted_test_xgb),s=0.1)
```

```
[ ]: <matplotlib.collections.PathCollection at 0x7f401646de90>
```



#####model prediction on test set:

Because we do not have test set for the submission version of logerror prediction, which was the scoring metric for Kaggle competition. Instead, below, we analyze the prediction on test set.


```
[ ]: y_pred = model_xgb.predict(dvalid)
```

```
[ ]: describe_arr(y_pred)
```

```
mean: 0.0095409425
median: 0.009667605
min: -0.109619394
max: 0.23125656
```

```
[ ]: plt.scatter(range(len(y_pred)), np.sort(y_pred),s=0.2)
plt.title("scatter plot for prediction")
```

```
[ ]: Text(0.5, 1.0, 'scatter plot for prediction')
```



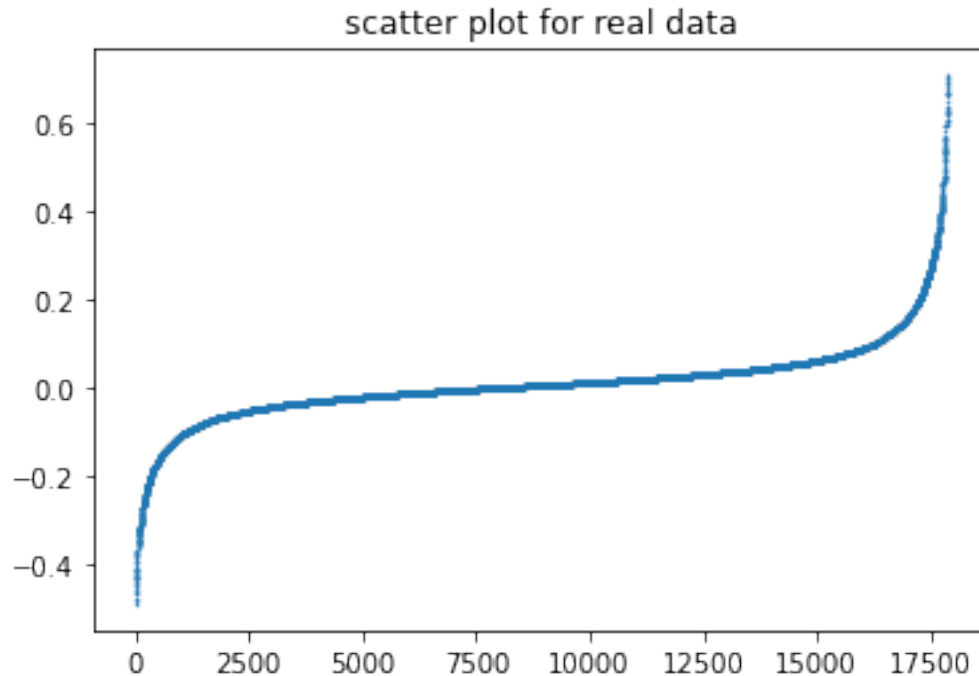
####real test data:

```
[ ]: describe_arr(yvalid)
```

```
mean: 0.011792033120734027
median: 0.007
min: -0.4878
max: 0.7066
```

```
[ ]: plt.scatter(range(len(yvalid)), np.sort(yvalid),s=0.2)
plt.title("scatter plot for real data")
```

```
[ ]: Text(0.5, 1.0, 'scatter plot for real data')
```



#####classification of logerror

a set a threshold of 0.1.

For logerror within 0.1 to -0.1 interval, we say we have a good estimate for this estate and represent it as 1. If the logerror is larger, it is replaced with 0, showing that it's an inaccurate estimate

```
[ ]: import copy
def classify_log(array):
    res = copy.deepcopy(array)
    for i in range(len(res)):
        if abs(res[i])<=0.1: res[i]=1
        else: res[i]=0
    return res
```

real test data:

```
[ ]: ytest_class = classify_log(yvalid)
```

```
[ ]: unique, counts = np.unique(ytest_class, return_counts=True)
dict(zip(unique, counts))
```

```
[ ]: {0.0: 2715, 1.0: 15159}
```

predicted data:

```
[ ]: ypred_class = classify_log(y_pred)
```

```
[ ]: unique, counts = np.unique(ypred_class, return_counts=True)
dict(zip(unique, counts))
```


```
[ ]: {0.0: 14, 1.0: 17860}
```

3.0.1 performance and fairness

```
[ ]: import sklearn.metrics as skm
print("accuracy:", skm.accuracy_score(ytest_class, ypred_class))
print("AUC:", skm.roc_auc_score(ytest_class, ypred_class))
print('precision:', skm.precision_score(ytest_class, ypred_class))
print('recall:', skm.recall_score(ytest_class, ypred_class))
print('average_precision:', skm.average_precision_score(ytest_class,
→ypred_class))
```

```
accuracy: 0.8481033903994629
AUC: 0.5010582484959611
precision: 0.8483762597984322
recall: 0.999538228115311
average_precision: 0.8483761337950155
```

```
[ ]: cm = skm.confusion_matrix(ytest_class, ypred_class)
skm.ConfusionMatrixDisplay(confusion_matrix=cm).plot()
plt.show()
```




```
[ ]: cm
```

```
[ ]: array([[ 7, 2708],  
          [ 7, 15152]])
```

```
[ ]: true_positive=cm[1,1]  
     false_positive=cm[0,1]  
     true_negative=cm[0,0]  
     false_negative=cm[1,0]
```

```
[ ]: print("false positive rate is: {:.4f}".format(false_positive/  
     ↪(false_positive+true_negative)))  
     print("false negative rate is: {:.4f}".format(false_negative/  
     ↪(false_negative+true_positive)))
```

```
false positive rate is: 0.9974
```

```
false negative rate is: 0.0005
```

3.0.2 attempt to analyze subpopulation accuracy

Because the ADS converted dataframes into numerical values when splitting training and test dataset. it is hard to reconstruct test set dataframe and reassociate predictions with features.

In addition, many features like buildingclasstypeid, storytypeid, architecturalstyletypeid obtains about 99% of null data. And features like regionidcounty, regionidcity, and regionidneighborhood do not have human-interpretable dictionary on Kaggle.

Therefore, although we originally planned to evaluate ADS's accuracy across subpopulations, we failed to analyze it due to ADS's encoding and preprocessing methods as well as missing documentations.

In the cells below, we attempted to analyze performance of subgroups based on "propertylanduse-typeid" by creating a fairlearn.metrics.MetricFrame taught in lab 4. However, in the process of coding, we realized we could not reconstruct predictions' associated features. Thus we failed to analyze fairness and subpopulation accuracy for this ADS due to above reasons.

```
[ ]: # !pip install --upgrade fairlearn==v0.6.0
```

```
[ ]: # from fairlearn.metrics import MetricFrame  
     # from fairlearn.metrics import selection_rate, false_negative_rate,  
     ↪false_positive_rate  
  
     # import functools  
     # import numpy as np
```

```
[ ]: # df_ytrain, df_yvalid = train_test_split(train, test_size=0.2, random_state=42)
```

```
[ ]: # def insensitive_roc_auc(y_true, y_score):  
     #     #to handle subgroups with only one class.
```

```

#     if sum(y_true)!=len(y_true) and sum(y_true)!=0:
#         return skm.roc_auc_score(y_true, y_score)
#     else:
#         return np.nan

# def samplesize(y_true, y_score):
#     return len(y_true)

# #Metrics
# metric_fns = {'samplesize': samplesize,
#               'selection_rate': selection_rate, # i.e., the percentage of
#               ↳the population which have '1' as their label
#               'FNR': false_negative_rate,
#               'FPR': false_positive_rate,
#               'accuracy': skm.accuracy_score,
#               'average_precision': skm.average_precision_score,
#               'roc_auc_score': insensitive_roc_auc
#               }

# grouped_on_race = MetricFrame(metric_fns,
#                                ytest_class, ypred_class,
#                                ↳
#                                ↳sensitive_features=df_yvalid['propertylandusetypeid'])

```

3.0.3 XGB factor importance

```

[ ]: weight = model_xgb.get_score(importance_type='weight')

[ ]: # removed columns from training set in pre-processing
removed_col = ['parcelid', 'logerror', 'transactiondate', 'propertyzoningdesc',
↳'propertycountylandusecode' ]

[ ]: # match unreadable coded names to actual feature names
feat_code = model_xgb.feature_names
feats_names = df_train.drop(removed_col, axis=1).columns


[ ]: code_to_names = {}
for key, value in zip(feat_code, feats_names):
    code_to_names[key] = value

[ ]: weight_dic = { code_to_names[k]: v for k, v in weight.items() }

[ ]: # sort dictionary from largest weight to lowest
import operator
sorted_weight = sorted(weight_dic.items(), key=operator.itemgetter(1),
↳reverse=False)

```

```
[ ]: x, y = zip(*sorted_weight)
      figs, axes = plt.subplots(figsize=(10,20))
      plt.barh(x,y)
      plt.show()
```



4 Bibliography

Zillow Zestimate ADS: <https://www.kaggle.com/c/zillow-prize-1/overview>

ADS Intended Goal and Specifications: <https://www.kaggle.com/c/zillow-prize-1/data>